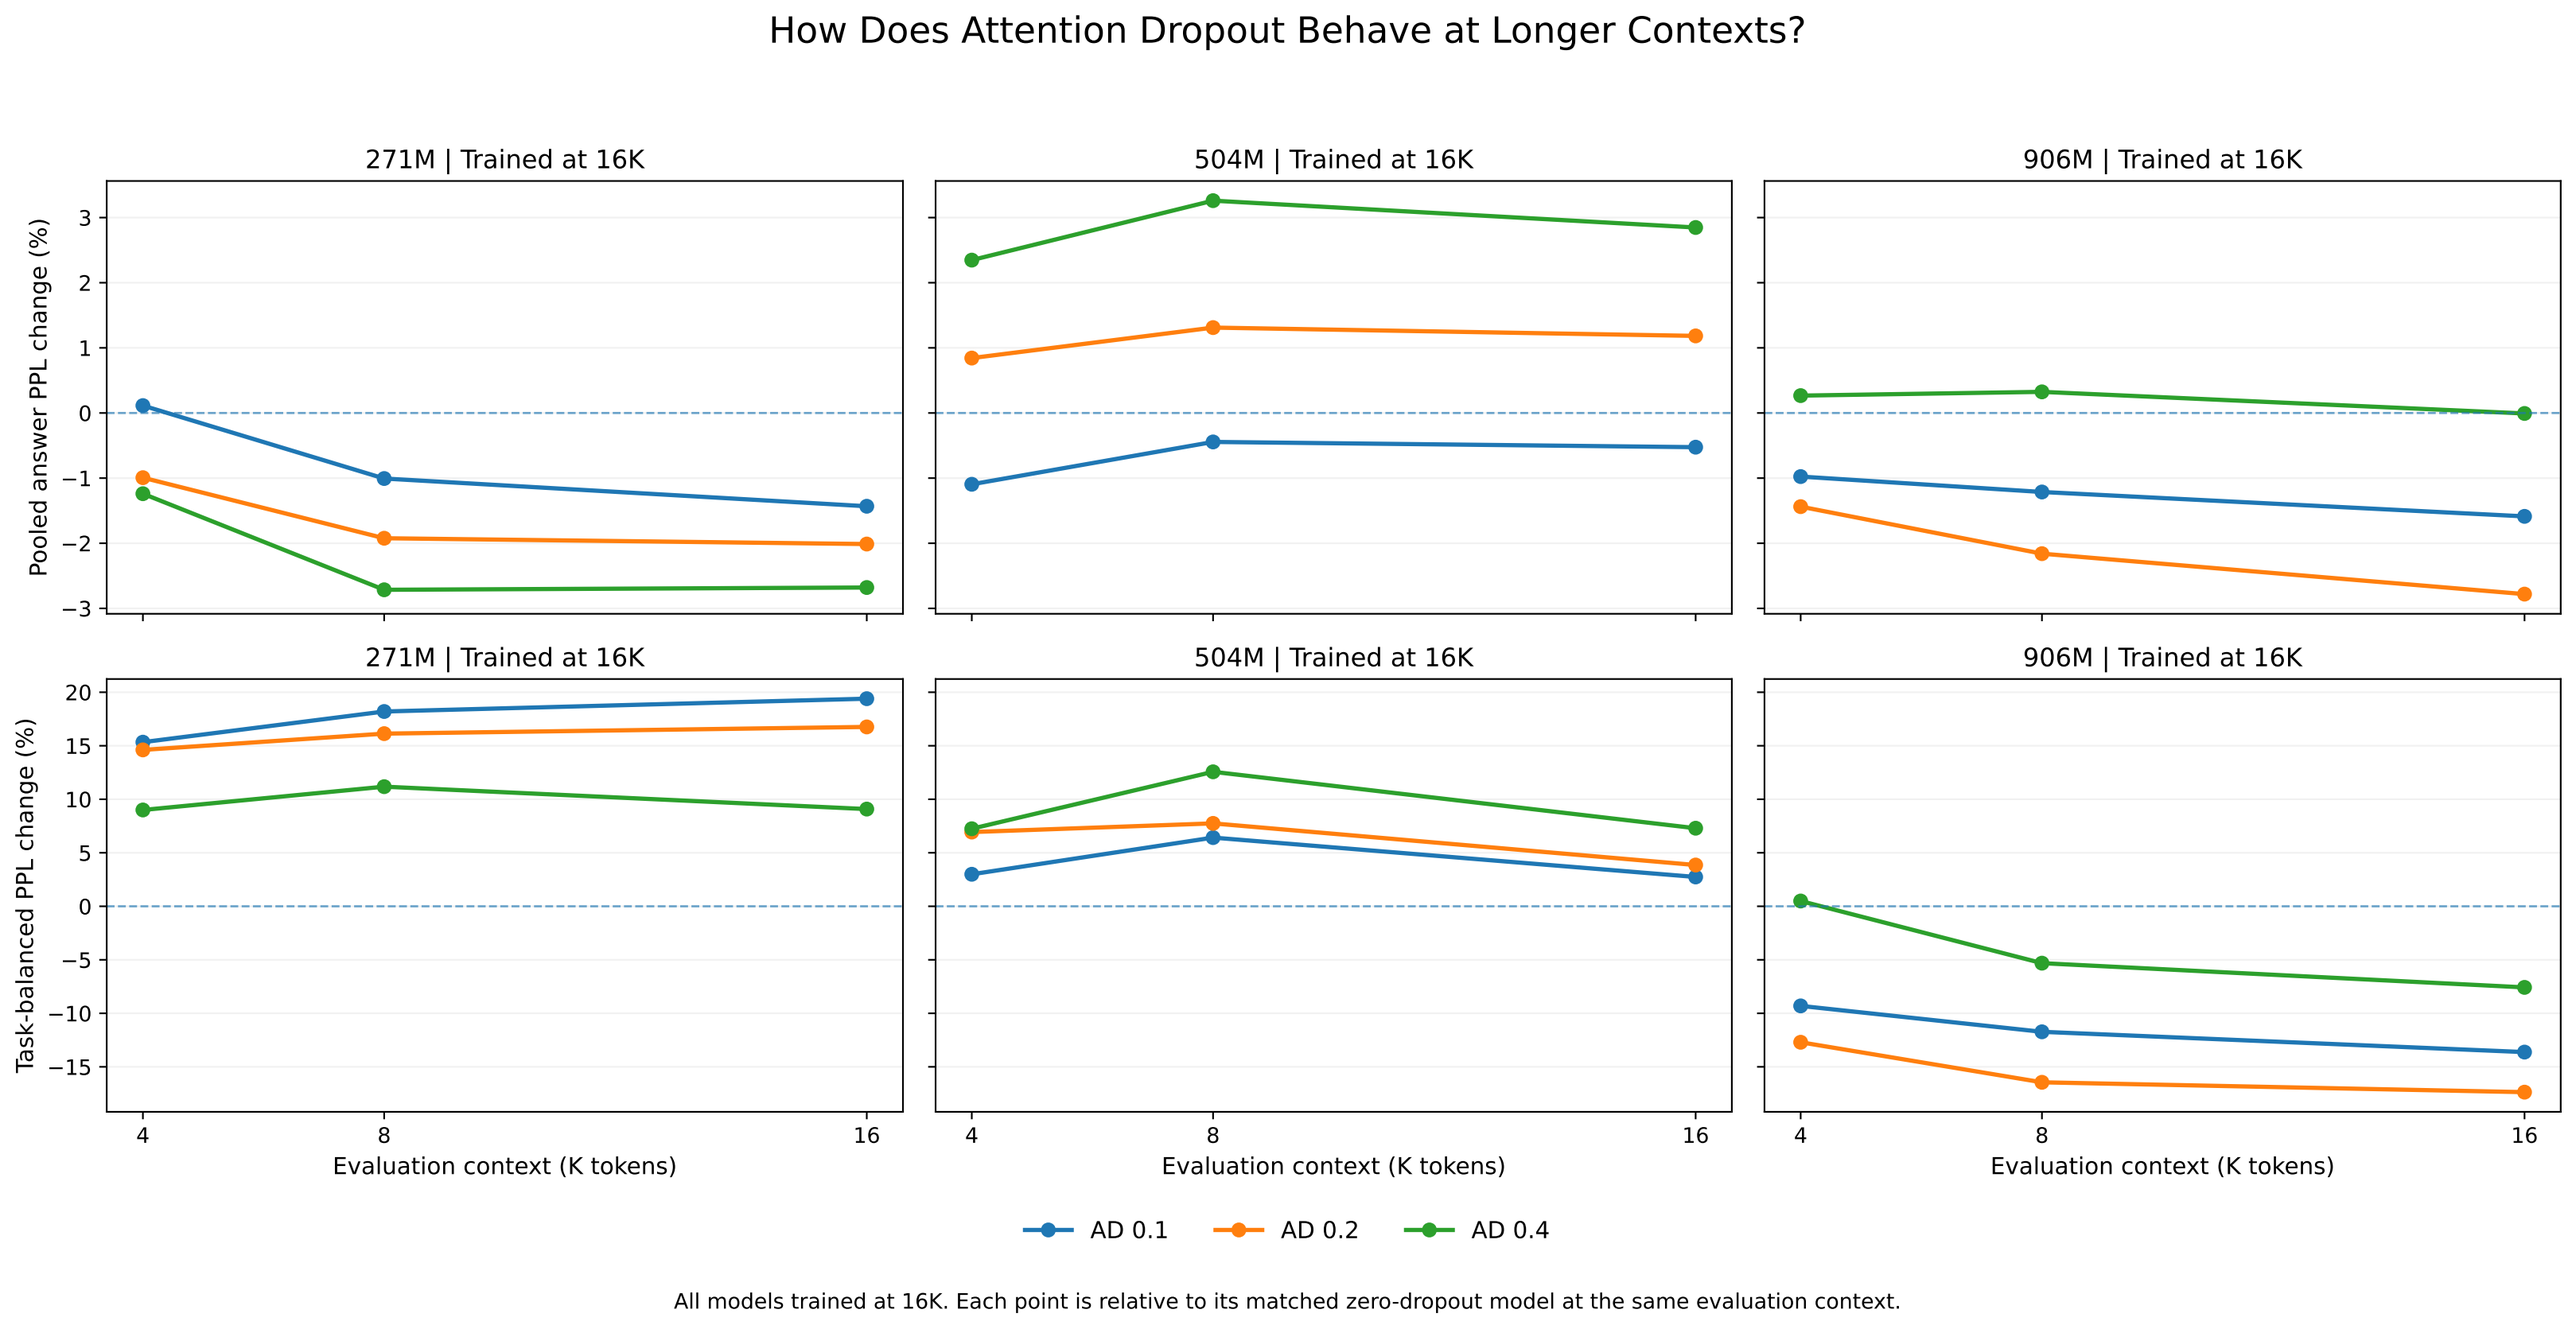
How Does Attention Dropout Behave at Longer Contexts?
271M | Trained at 16K
504M | Trained at 16K
906M | Trained at 16K
Pooled answer PPL change (%)
3
2
1
0
−1
−2
−3
271M | Trained at 16K
504M | Trained at 16K
906M | Trained at 16K
Task-balanced PPL change (%)
20
15
10
5
0
−5
−10
−15
4
8
16
Evaluation context (K tokens)
Evaluation context (K tokens)
Evaluation context (K tokens)
AD 0.1
AD 0.2
AD 0.4
All models trained at 16K. Each point is relative to its matched zero-dropout model at the same evaluation context.